

Compute Evolved Week

Building Microservices with the 12 Factor App Pattern on AWS

Nathan Peck

Developer Advocate, Container Services



12 Factor App
Principles



+

Microservice
Principles




=

Great, Scalable
Architecture





12 Factor Application: Codebase

Code





Code



Version Control



Code




Version Control




git

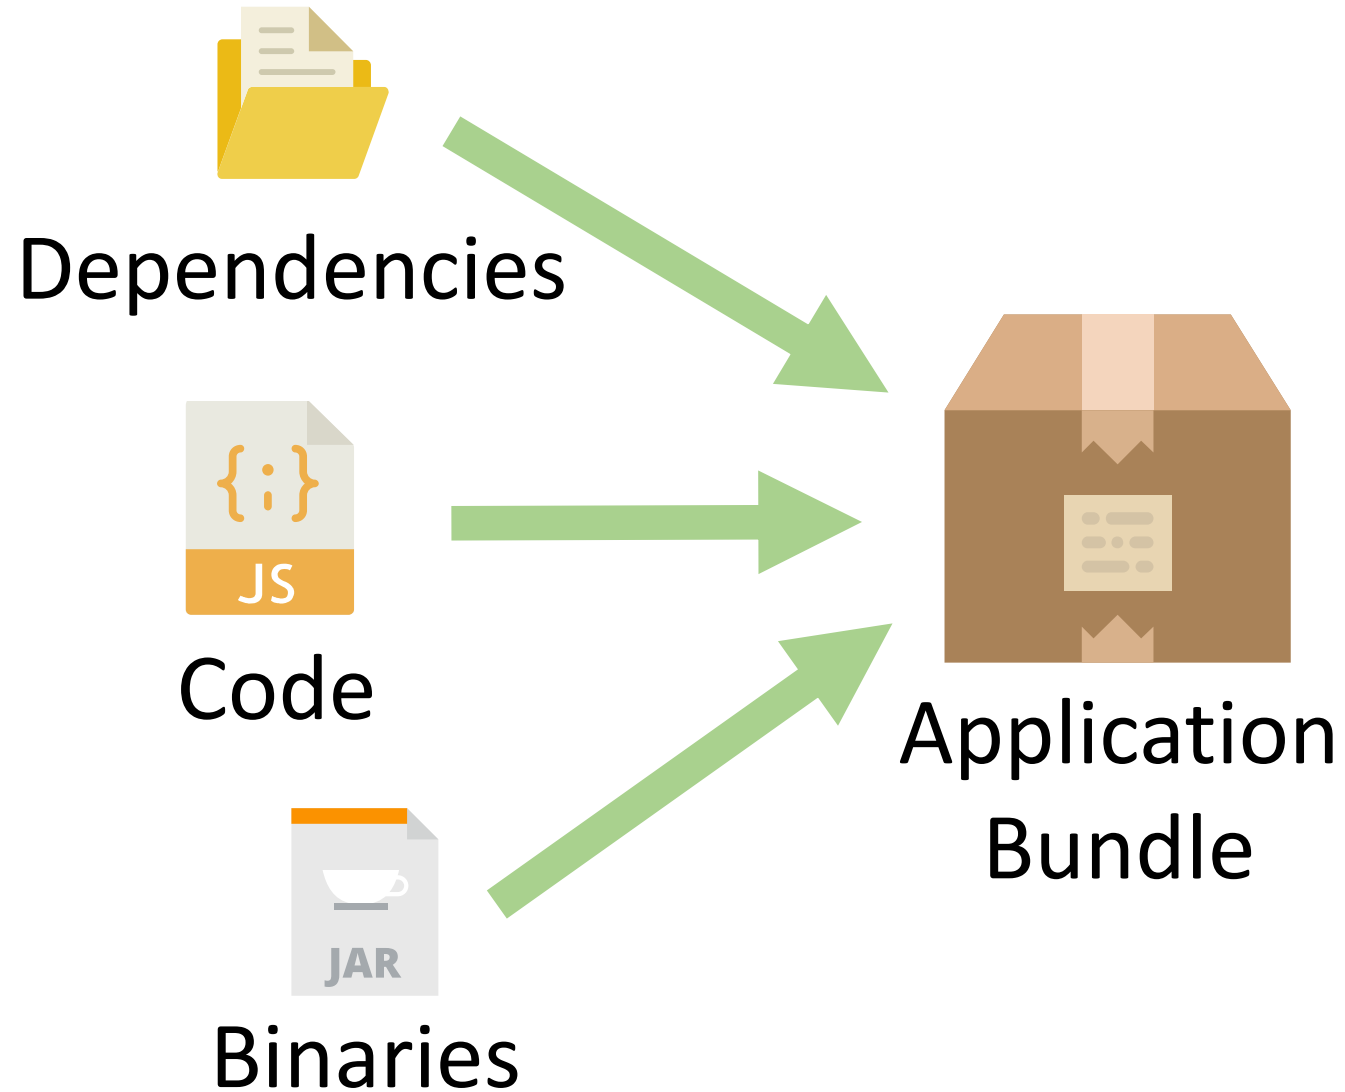


Deployed Version





12 Factor Application: Dependencies



Dependency Declaration: Node.js

package.json

```
{
  "dependencies": {
    "async": "2.1.4",
    "express": "4.16.2",
    "express-bearer-token": "2.1.0",
    "body-parser": "1.18.2",
    "jwt-simple": "0.5.1",
    "lodash": "4.17.4",
    "morgan": "1.7.0",
    "request": "2.81.0"
  }
}
```

npm install

Dependency Declaration: Python

requirements.txt

```
django==1.6
bpython==0.12
django-braces==0.2.1
django-model-utils==1.1.0
logutils==0.3.3
South==0.7.6
requests==1.2.0
stripe==1.9.1
dj-database-url==0.2.1
django-oauth2-provider==0.2.4
django-rest-framework==2.3.1
```

pip install

Dependency Declaration: Ruby

Gemfile

```
source 'https://rubygems.org'


gem 'nokogiri'
gem 'rails', '3.0.0.beta3'
gem 'rack', '>=1.0'
gem 'thin', '~>1.1'
```

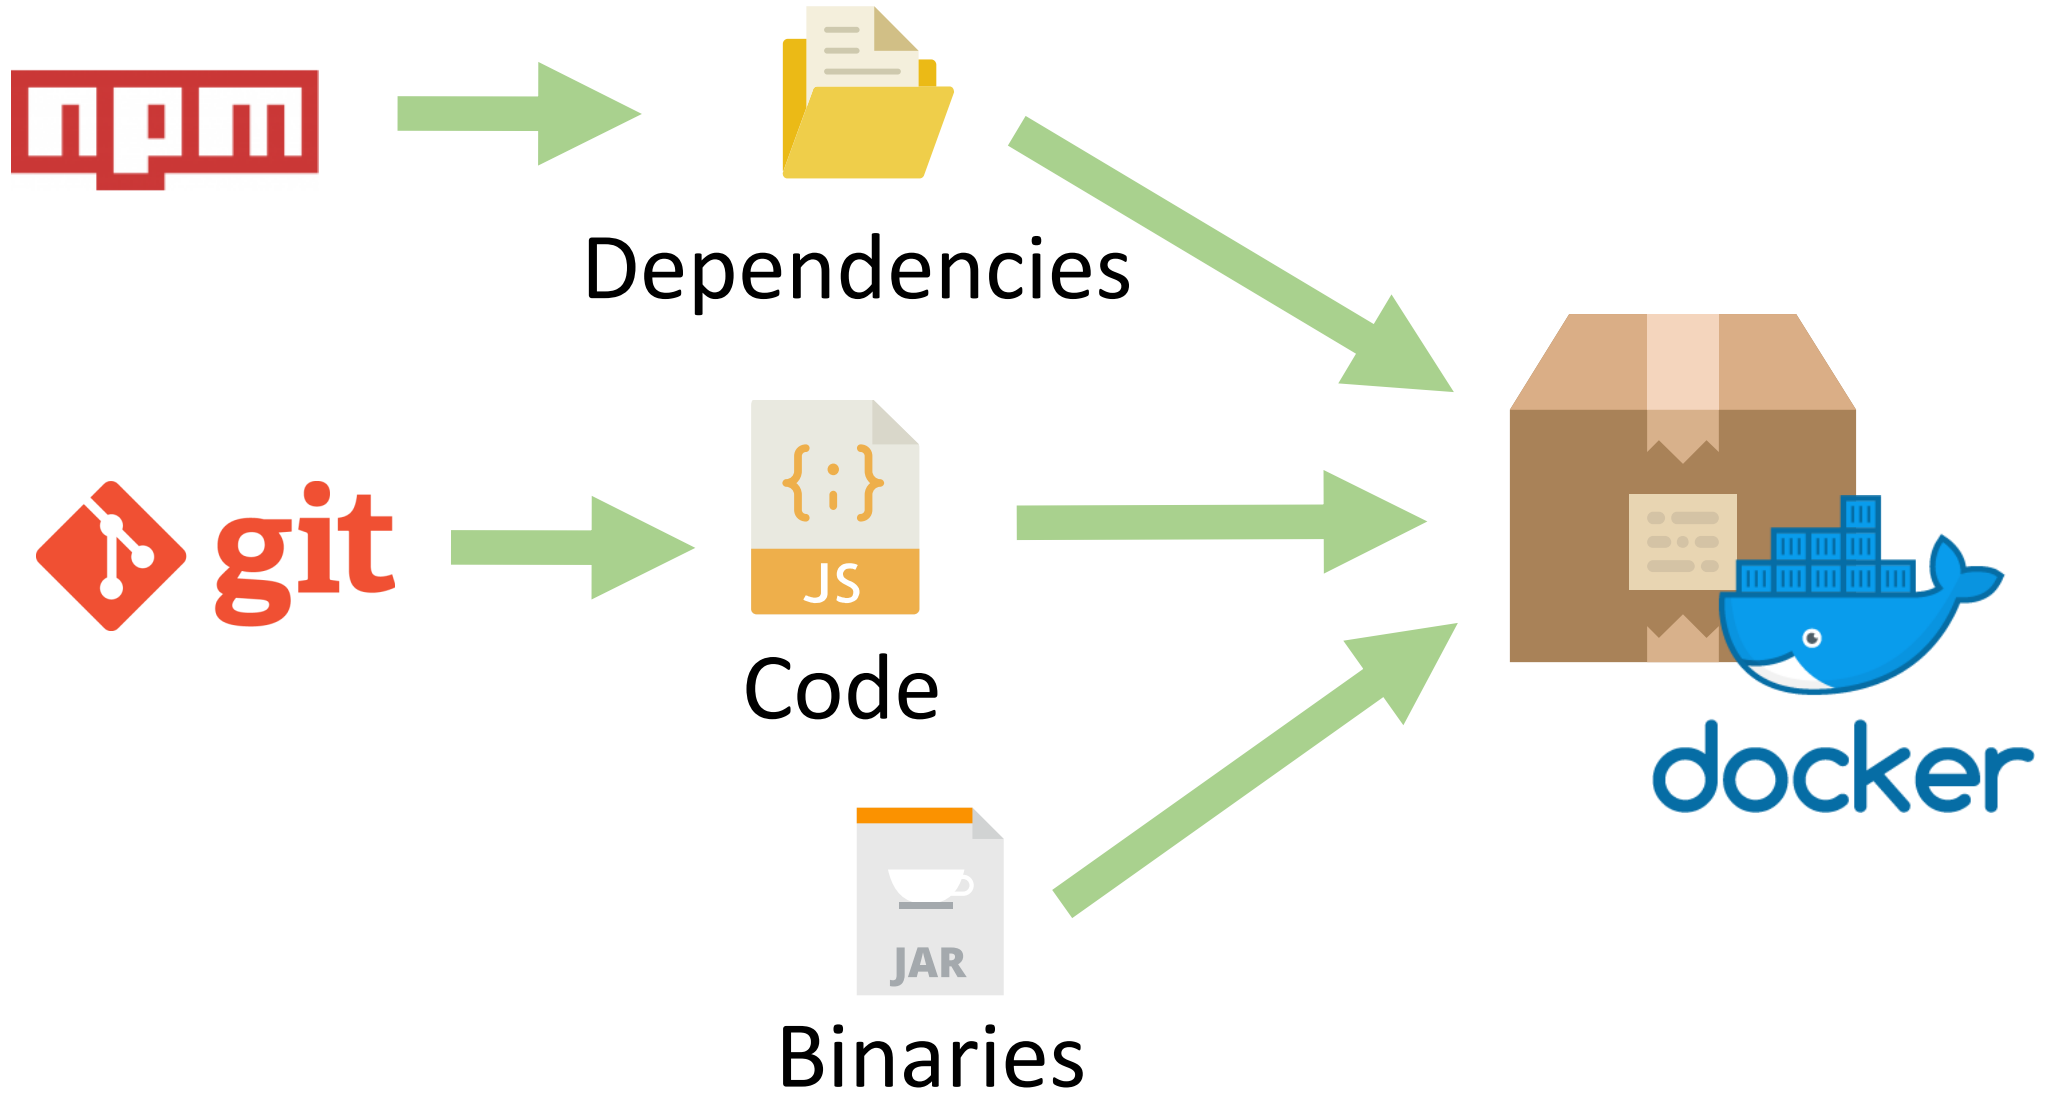
bundle install

Dependency Isolation

Never depend on the host to have your dependency.

Application deployments should carry all their dependencies with them.





Dependency Declaration & Isolation: Docker

Dockerfile

```
FROM mhart/alpine-node:8


RUN apk add --no-cache make gcc g++ python

WORKDIR /srv
ADD . .
RUN npm install


EXPOSE 3000
CMD ["node", "index.js"]
```

docker build

`docker run`



Development





Production

12 Factor Application: Config





Development
Configuration



Development




Production
Configuration



Production



Development
Configuration





Development

ANTIPATTERN




Production
Configuration




Production

Same container deployed to both environments.
Configuration is part of the environment on the host.



At runtime the container gets config from the environment.




Application code pulls from the environment

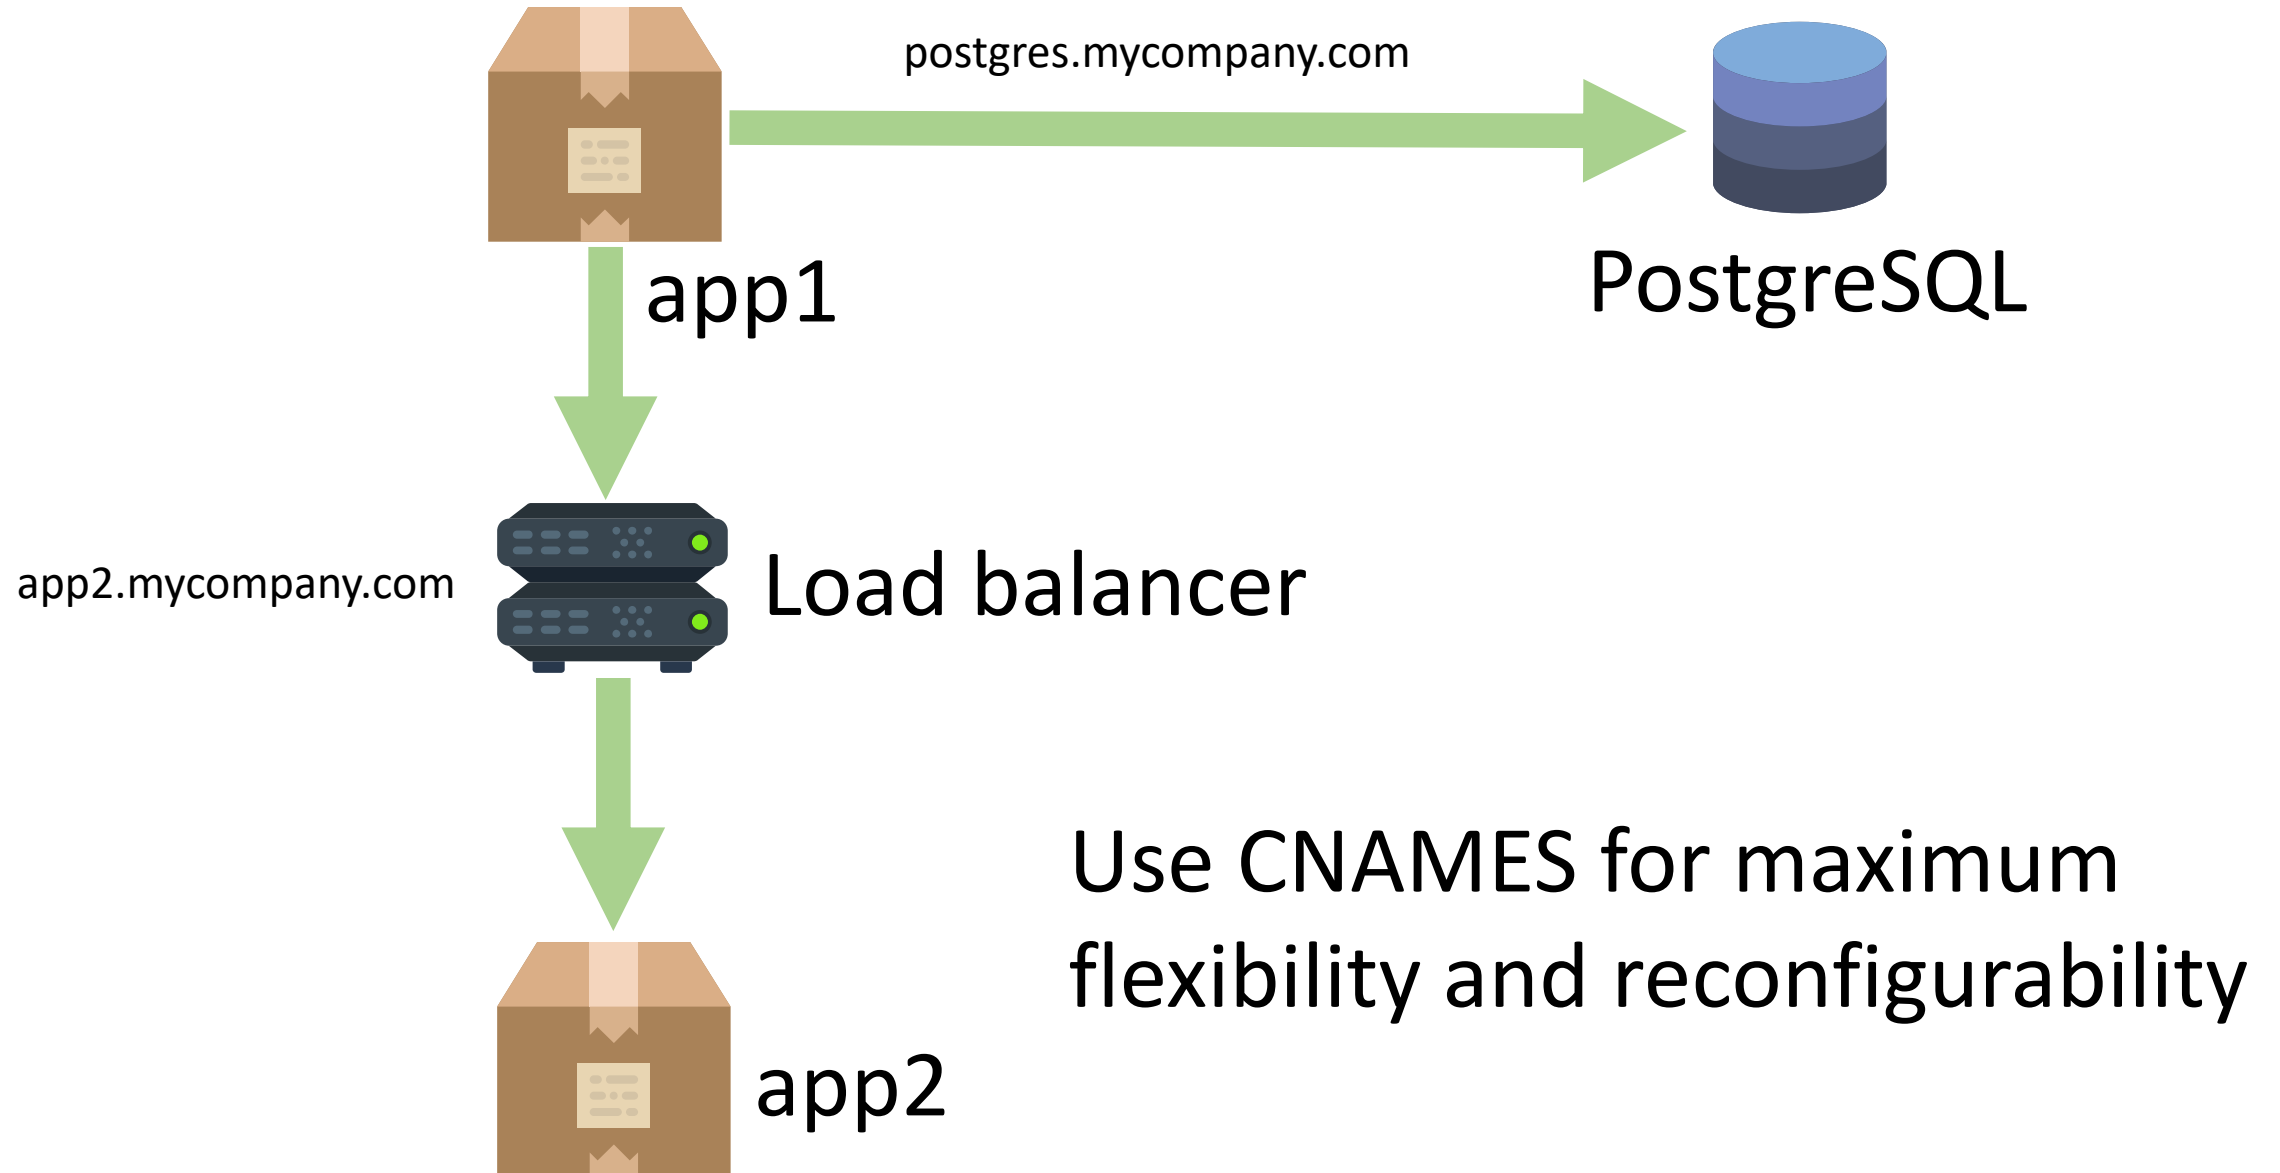
```
module.exports = {  
  DATABASE: process.env.DATABASE,  
  SECRET: process.env.SECRET  
};
```

Environment is customized when docker runs a container

```
docker run -e "DATABASE=mongodb://localhost:27017" -e "SECRET=default" myapp  
  
docker run -e "DATABASE=mongodb://db1.mycompany.com,db2.mycompany.com/  
production?replicaSet=production" -e "SECRET=hunter2" myapp
```

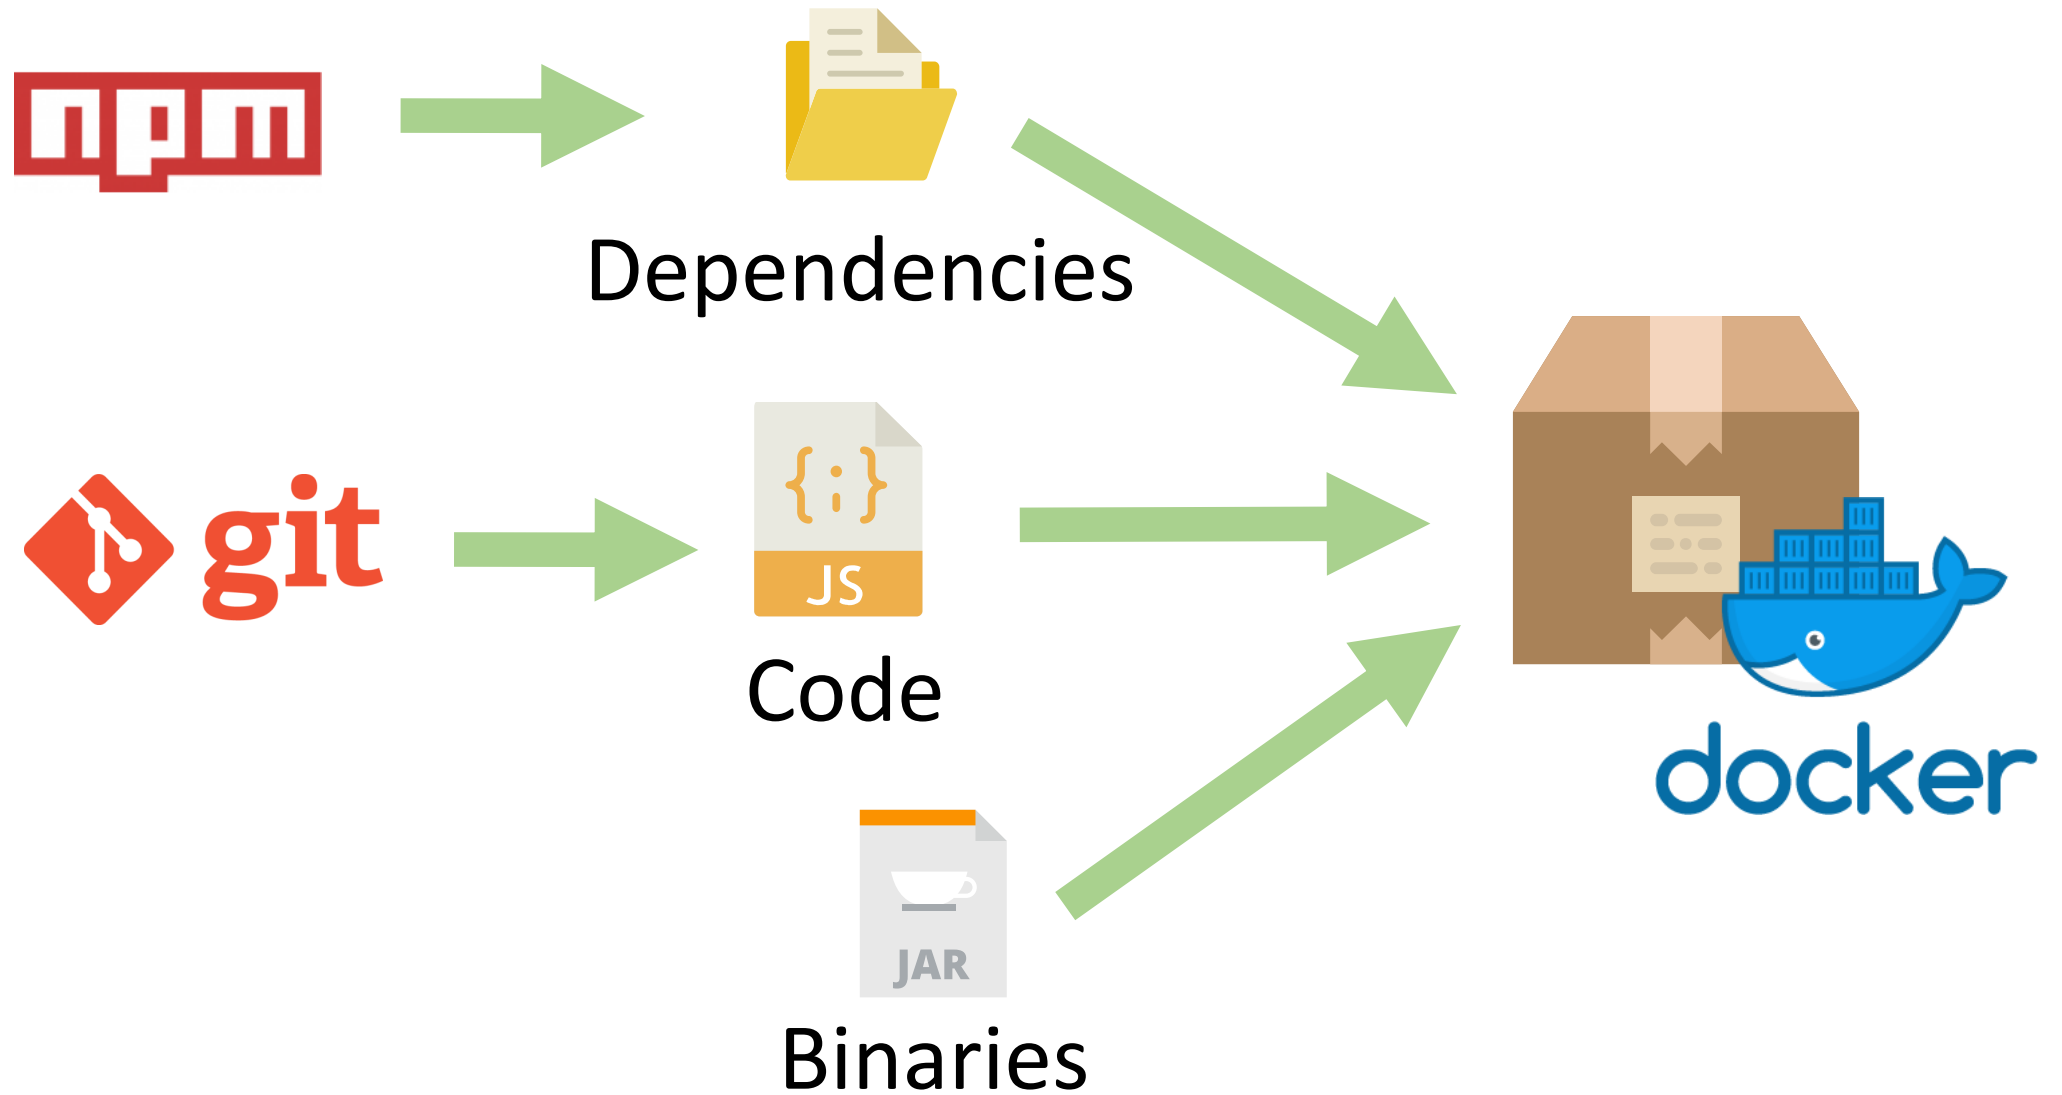
12 Factor Application: Backing Services






12 Factor Application: Build, Release, Run

Build



Release

Build Artifact



docker

+

Config



=

Release



docker

Amazon Elastic Container Service


```
# The task definition that describes how to run the docker container
TaskDefinition:
  Type: AWS::ECS::TaskDefinition
  Properties:
    Family: !Ref 'ServiceName'
    TaskRoleArn: !Ref TaskRole
    ContainerDefinitions:
      - Name: !Ref 'ServiceName'
        Cpu: 512
        Memory: 256
        Image: !Ref 'ImageUrl'
        Ulimits:
          - Name: nofile
            HardLimit: 65535
            SoftLimit: 65535
        PortMappings:
          - ContainerPort: 3000
            HostPort: 0
        Environment:
          - Name: 'UV_THREAD_POOL'
            Value: '15'
          - Name: 'NODE_ENV'
            Value: 'production'
```




Config

Run

Task Definition
Release v1.0.0




Task Definition
Release v1.0.1




12 Factor Application: Stateless Processes


Stateful container stores state in local disk or local memory.
Workload ends up tied to a specific host that has state data.




Stateful container stores state in local disk or local memory.
Workload ends up tied to a specific host that has state data.



ANTIPATTERN



docker




docker




Database

Durable store of truth.
MySQL, PostgreSQL,
MongoDB, DynamoDB





Cache

Fast, temporary
state store.
redis, memcached




12 Factor Application: Port Binding



Port 32456



docker





Port 32457




docker

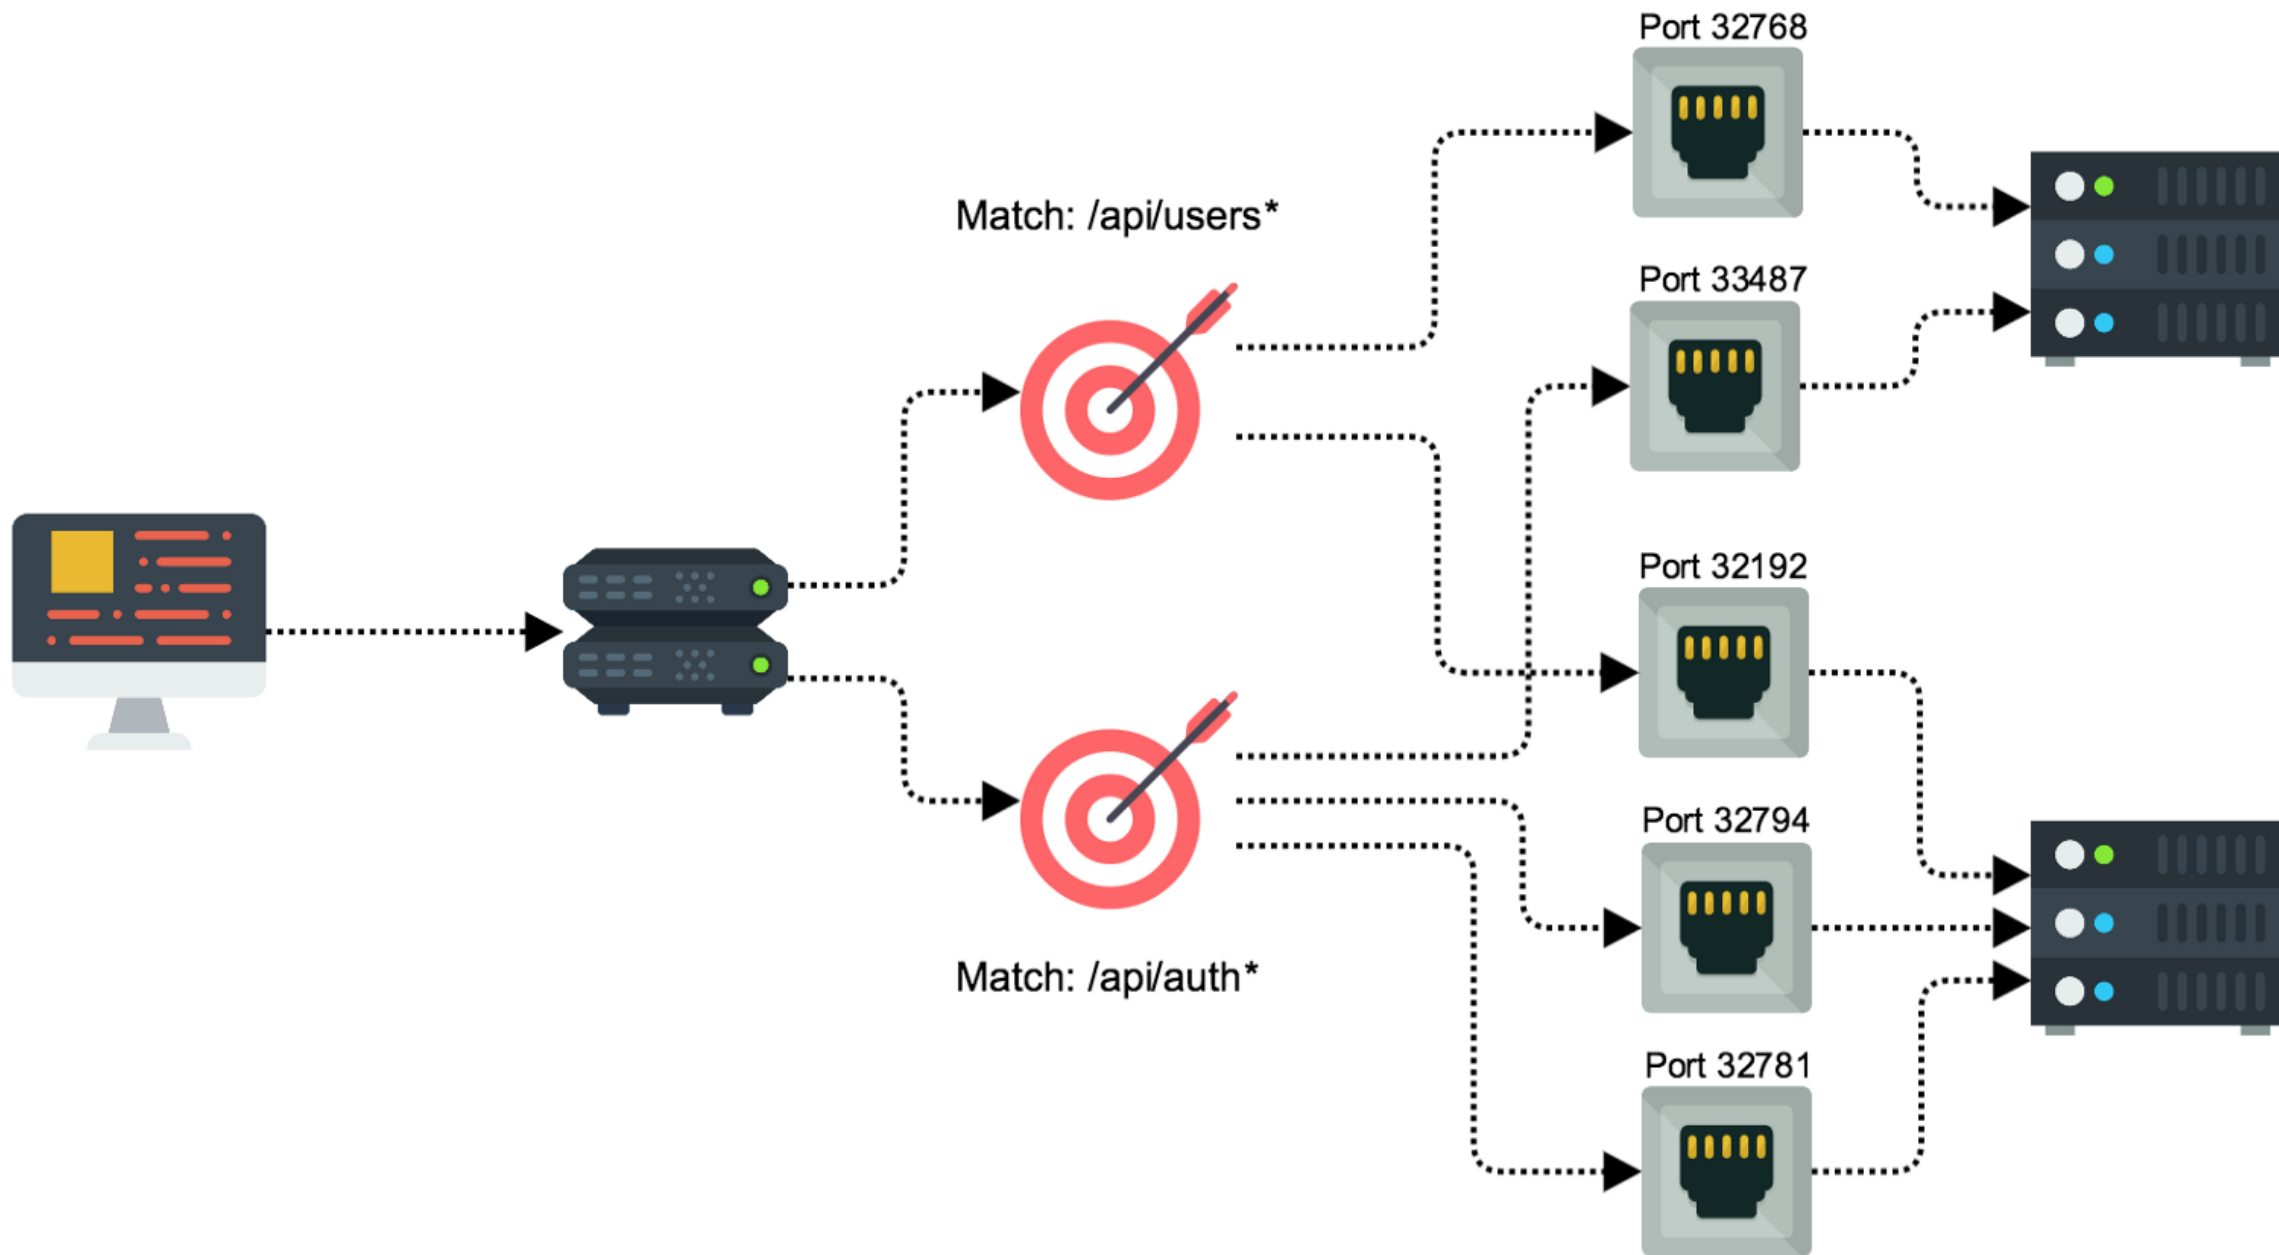


Port 32458





docker





12 Factor Application: Concurrency




API

Web

Worker


Web



API





Worker



Hosts





Processes




Hosts




Processes






Large Host = More
Concurrent Processes



Small Host =
Fewer Concurrent
Processes

12 Factor Application: Disposability



Fast Launch

Responsive

Graceful
Shutdown


Fast Launch



Minimize the startup time of processes:

- Scale up faster in response to spikes
- Ability to move processes to another host as needed
- Replace crashed processes faster


Responsive, Graceful Shutdown



Should respond to SIGTERM by shutting down gracefully

```
var server = app.listen(3000);  
  
console.log('Message service started');  
  
process.on('SIGTERM', function() {  
  console.log('Shutting down message service');  
  server.close();  
});
```


12 Factor Application: Dev/Prod Parity



Local




Application

Remote



Dev #1



Dev #2



docker



Staging / QA




Production




12 Factor Application: Logs

Treat logs as an event stream,
and keep the logic for routing
and processing logs separate
from the application itself.




Process



elastic

Process




Some logs get lost
if they haven't
fully flushed



elastic

ANTIPATTERN

Process




Some logs get lost
if they haven't
fully flushed



elastic

Processes

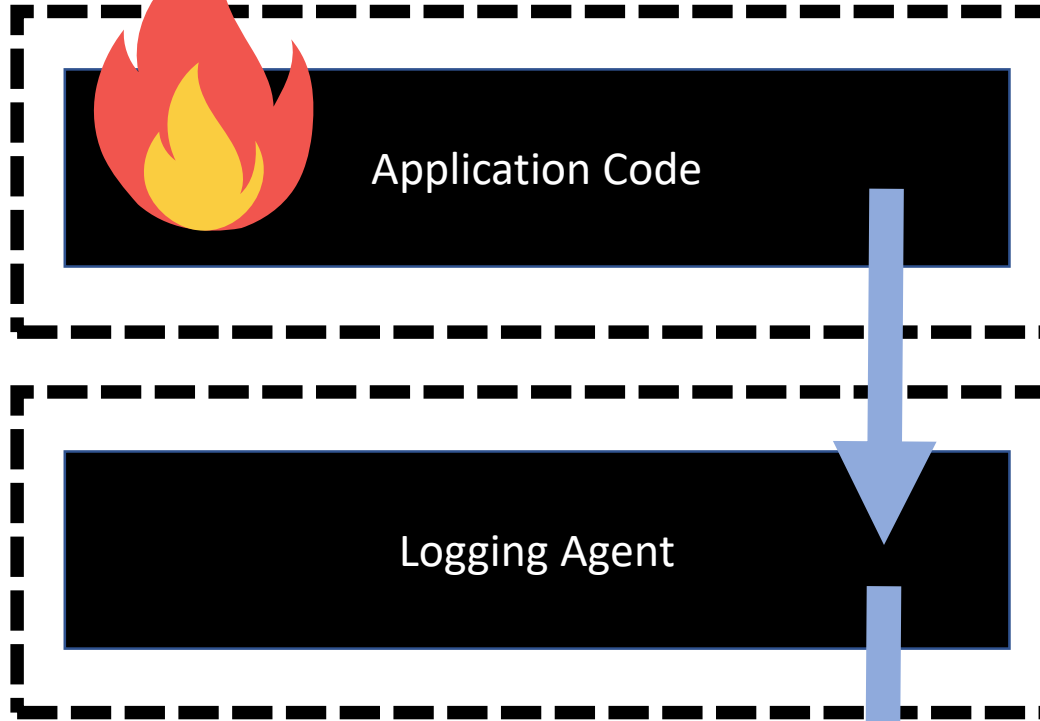


Logs go to an agent which handles exporting them off the host



elastic

Processes



Logs still reach agent, and still make it into ELK stack



elastic


Containerized code writes to stdout

```
var express = require('express');  
var app = express();  
var logger = require('morgan');  
  
app.use(logger('tiny'));
```

Docker connects container's stdout to a log driver


```
docker run --log-driver awslogs myapp  
  
docker run --log-driver fluentd myapp
```

12 Factor Application: Admin Processes




Admin / management processes are inevitable:

- Migrate database
- Repair some broken data
- Once a week move database records older than X to cold storage
- Every day email a report to this person




Run admin processes
just like other
processes.

Microservices: Componentization



Each component is a 12 factor application.




Codebase




Dependencies



Configuration



Port Binding




Stateless




Fast Launch



Graceful stop



Log stream




Concurrent

Microservices:
Organized around capabilities




Identify the capabilities of the platform



Each major capability of the platform becomes a component that is its own 12 factor app



Microservices: Decentralized Governance



Frontend

Game Client
(C++)

Website
(HTML, React)

Blog + CMS
(HTML, JS)

Backend

Game Server
(Java)

User Accounts
(Java)

React Renderer
(Node.js)

Microtransactions
(Node.js)

WordPress Backend
(PHP)

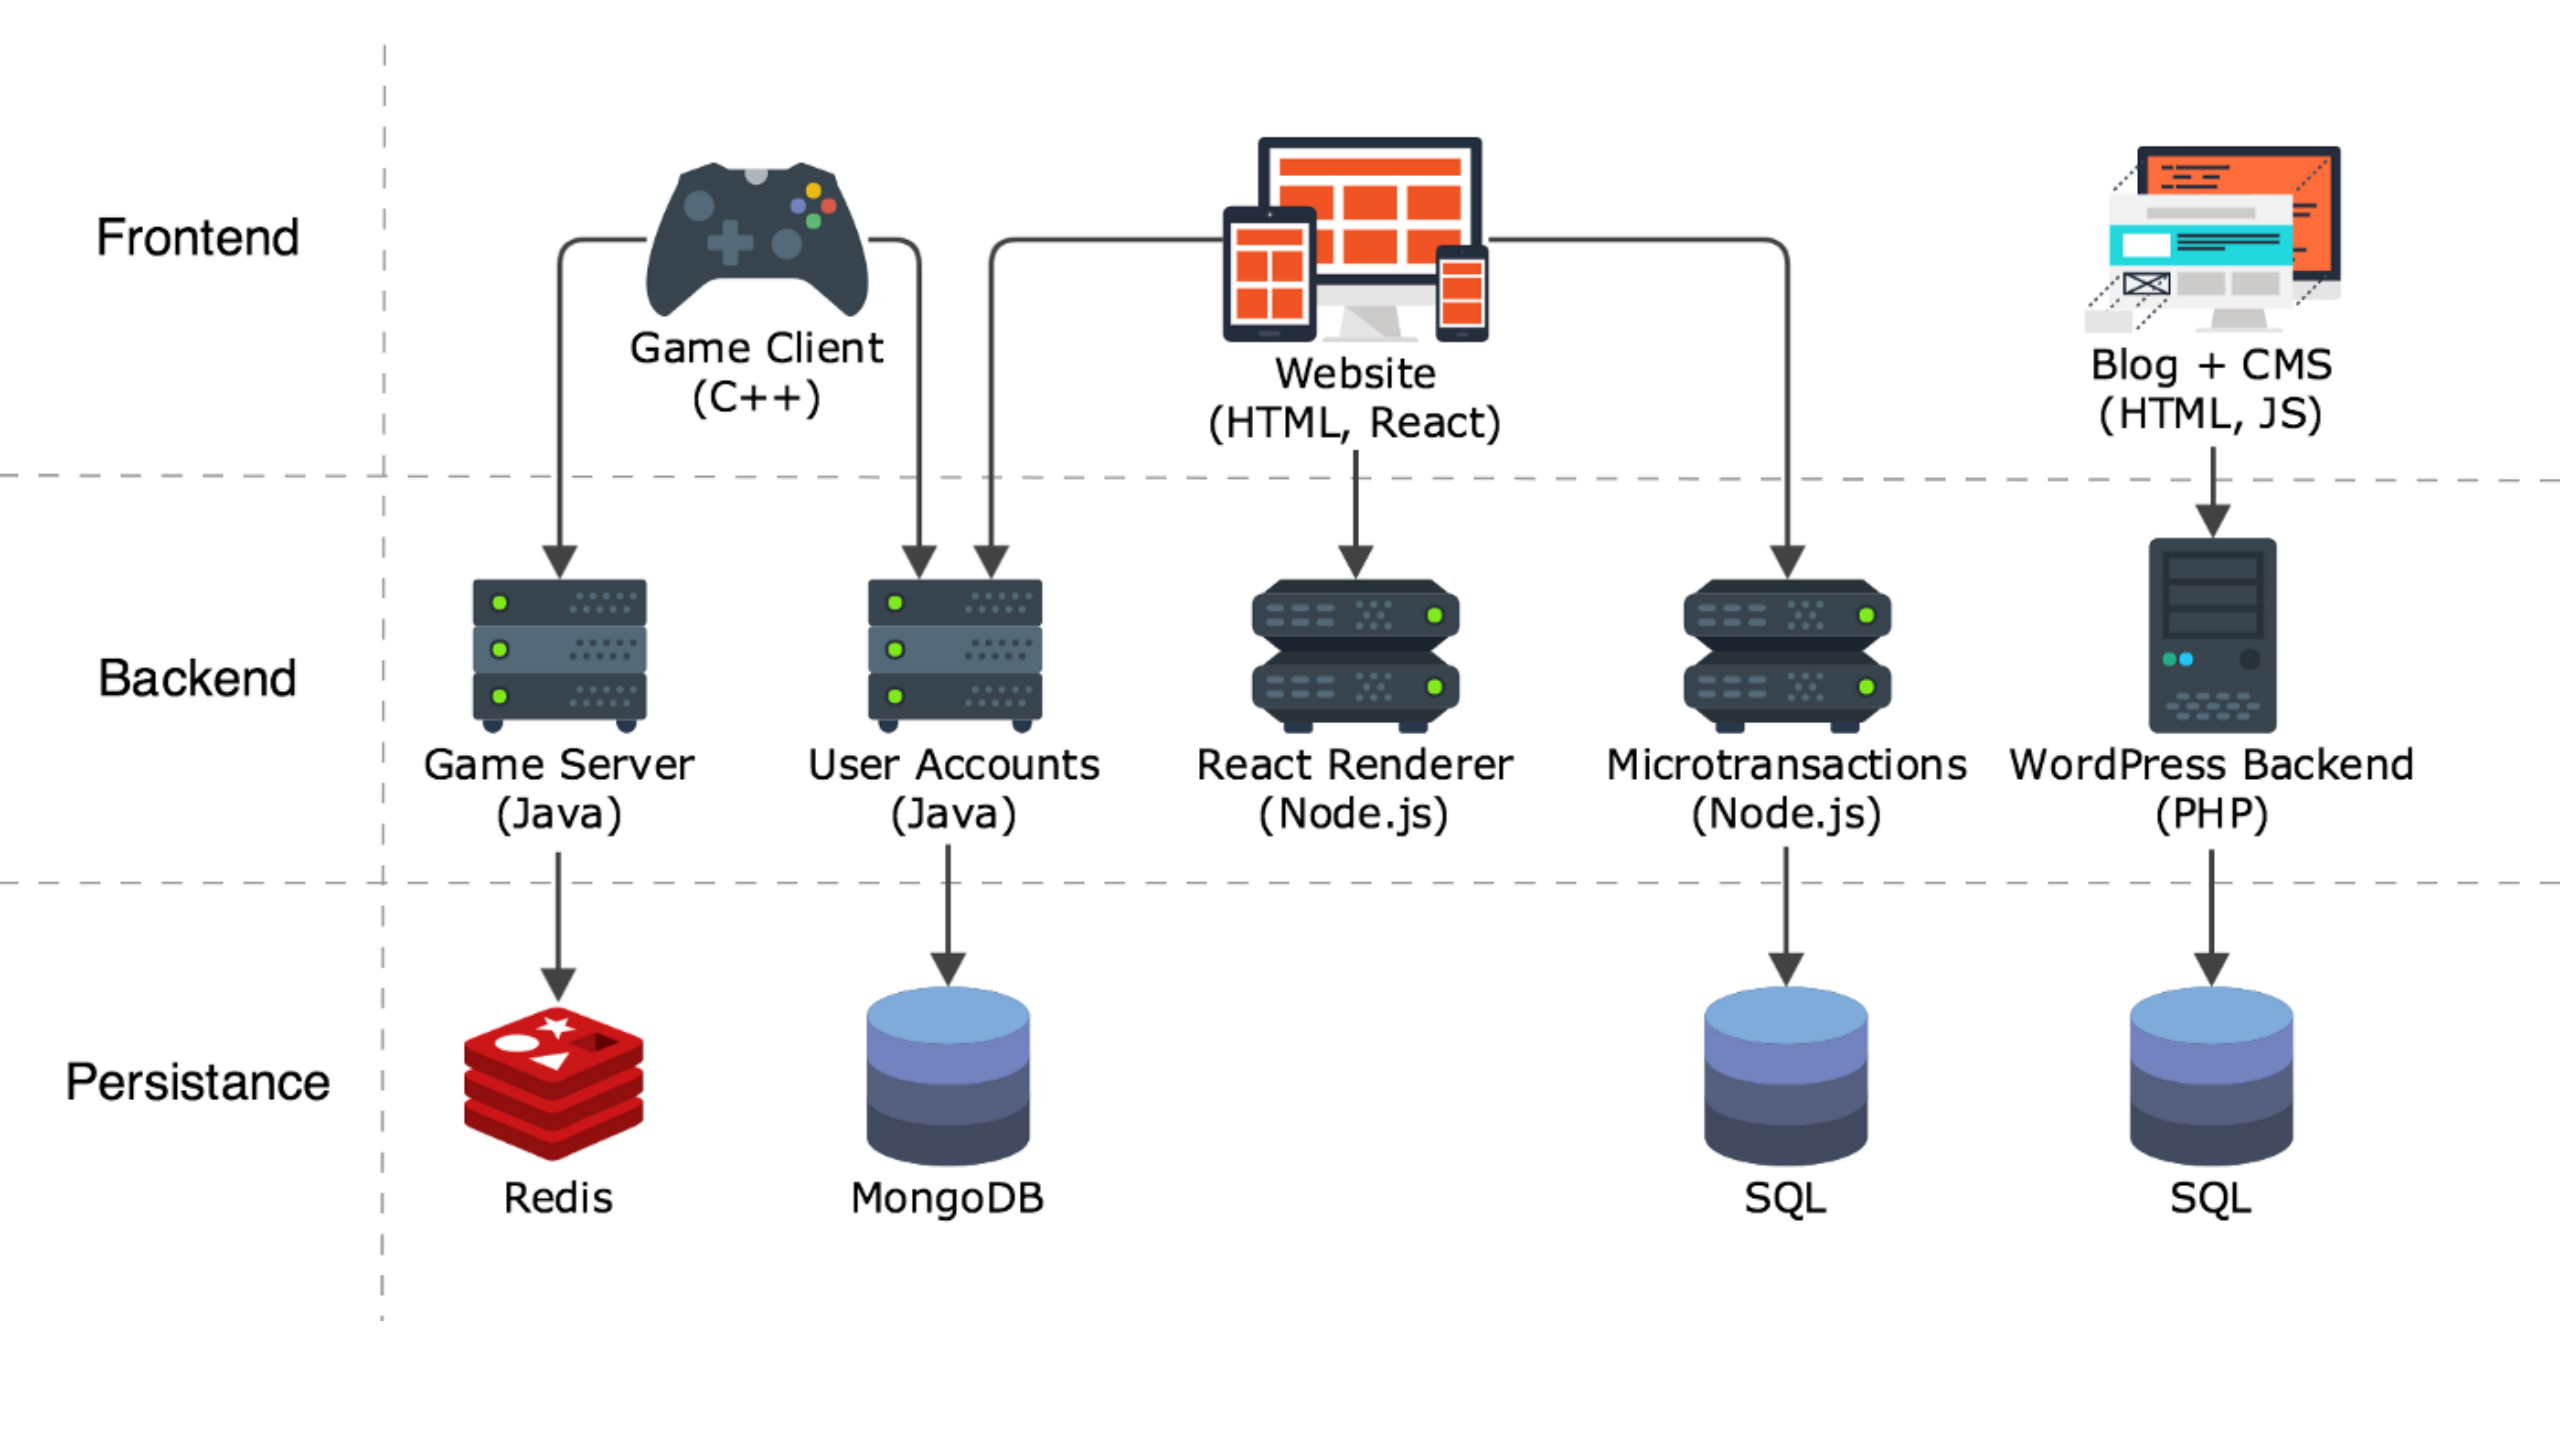
Persistence

Redis

MongoDB

SQL

SQL



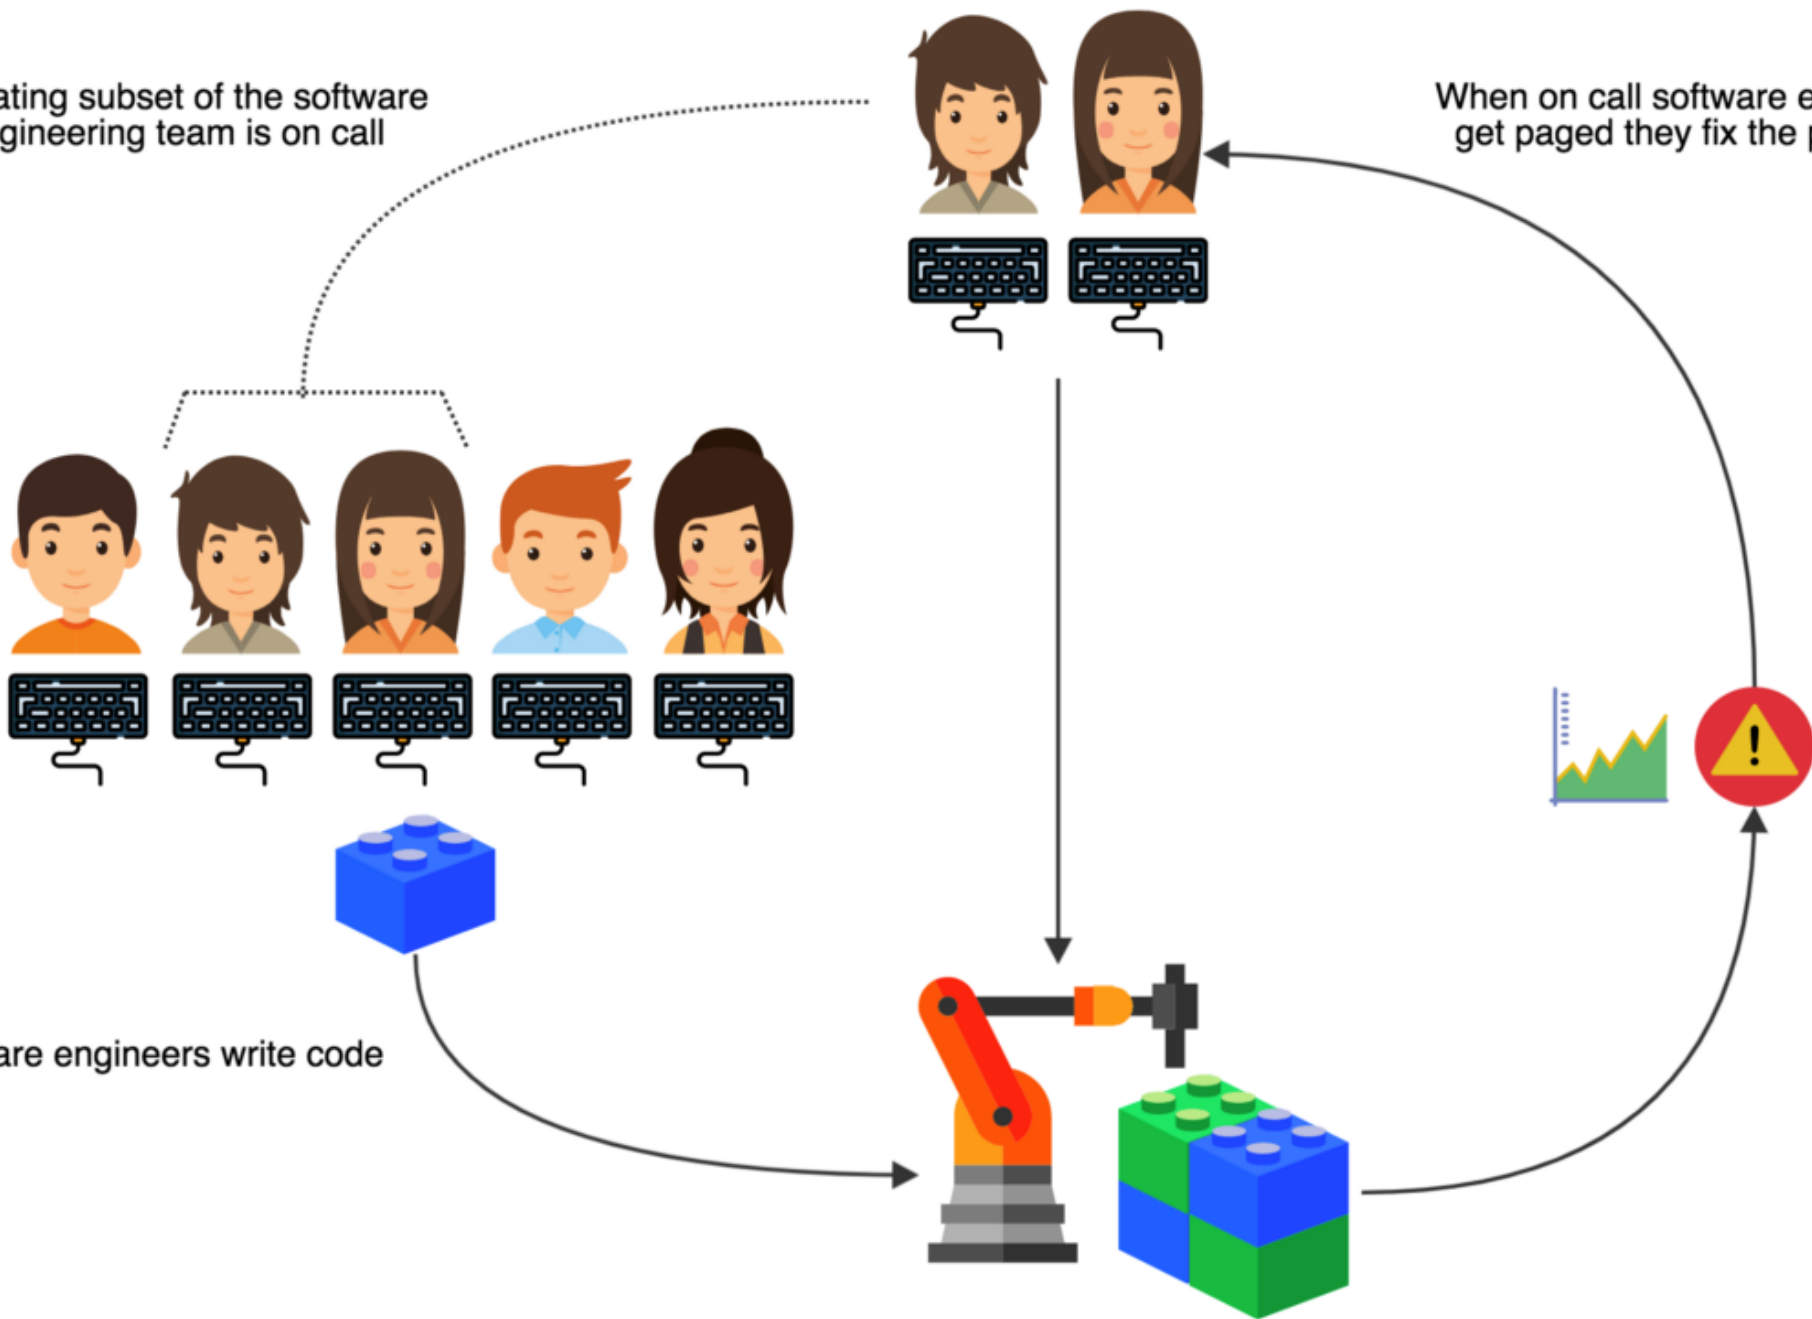
A rotating subset of the software engineering team is on call

When on call software engineers get paged they fix the problem

Software engineers write code



Automation deploys the code

Monitoring catches something bad happening in production and alerts are triggered




Microservices:
Products Not Projects


Products grow over time



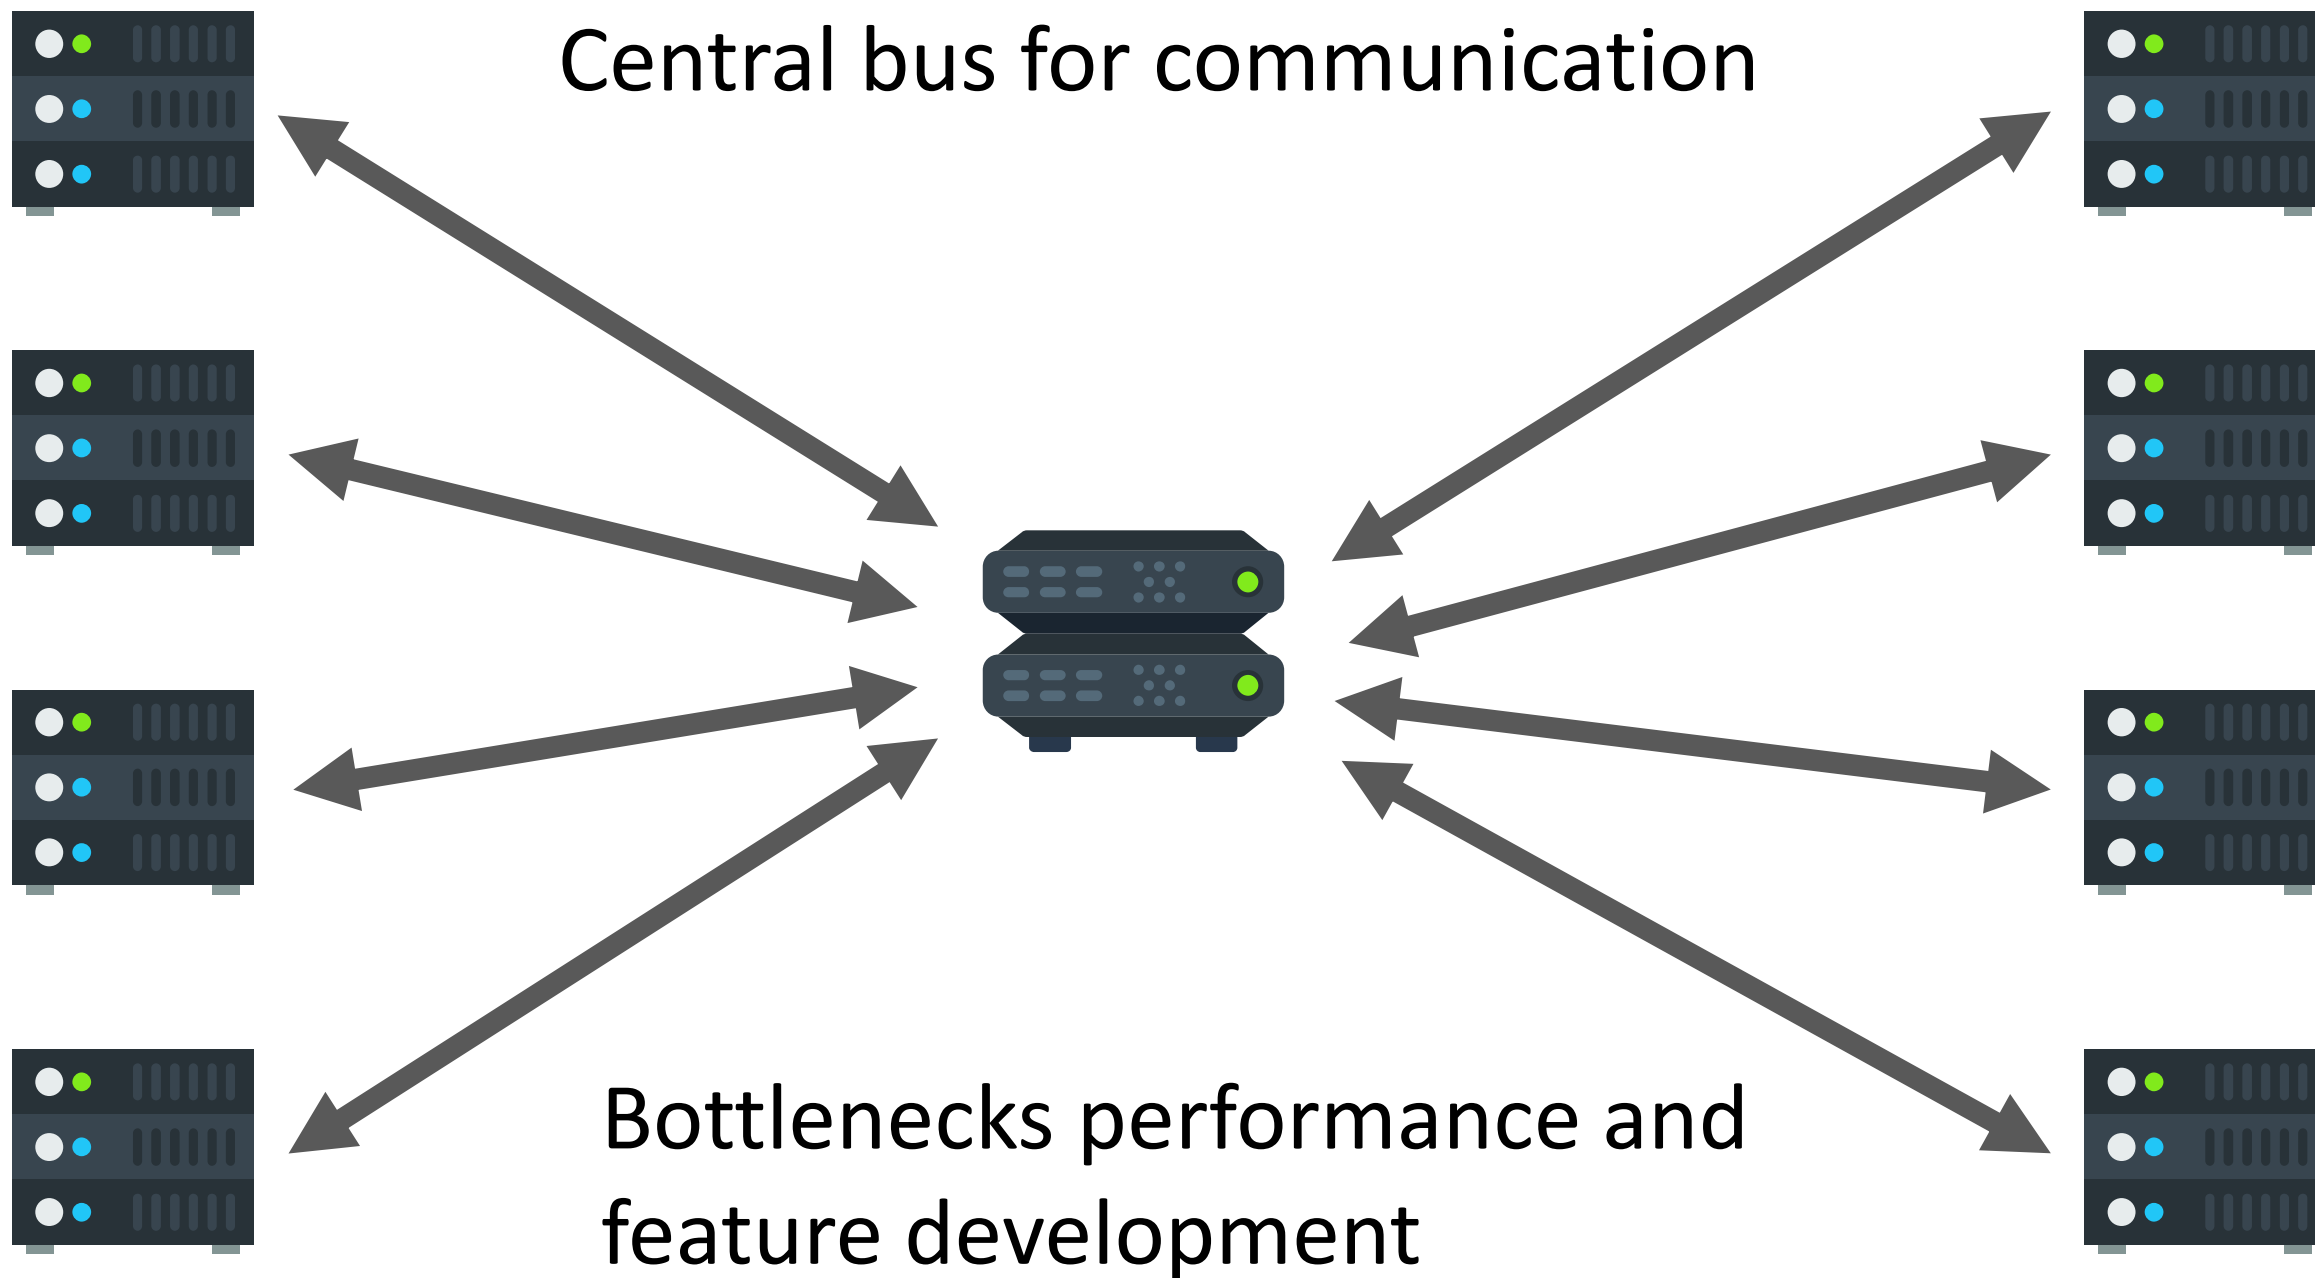
Microservices are an ecosystem of connected products that can be at different stages of growth



Don't create throwaway microservices that become unmaintained and break the ecosystem.






Microservices:
Smart endpoints, Dumb pipes




Central bus for communication

ANTIPATTERN


Bottlenecks performance and
feature development



Decentralization



Example: User Signup








Event Producers

Event Topics

Subscriptions

Queues

Event Consumers



Amazon Managed Service for Microservice Communication



Amazon Simple Notification Service (SNS)




Amazon Simple Queue Service (SQS)

Amazon MQ

Managed message broker service for Apache ActiveMQ



Microservices: Infrastructure Automation

Automate the container build process:



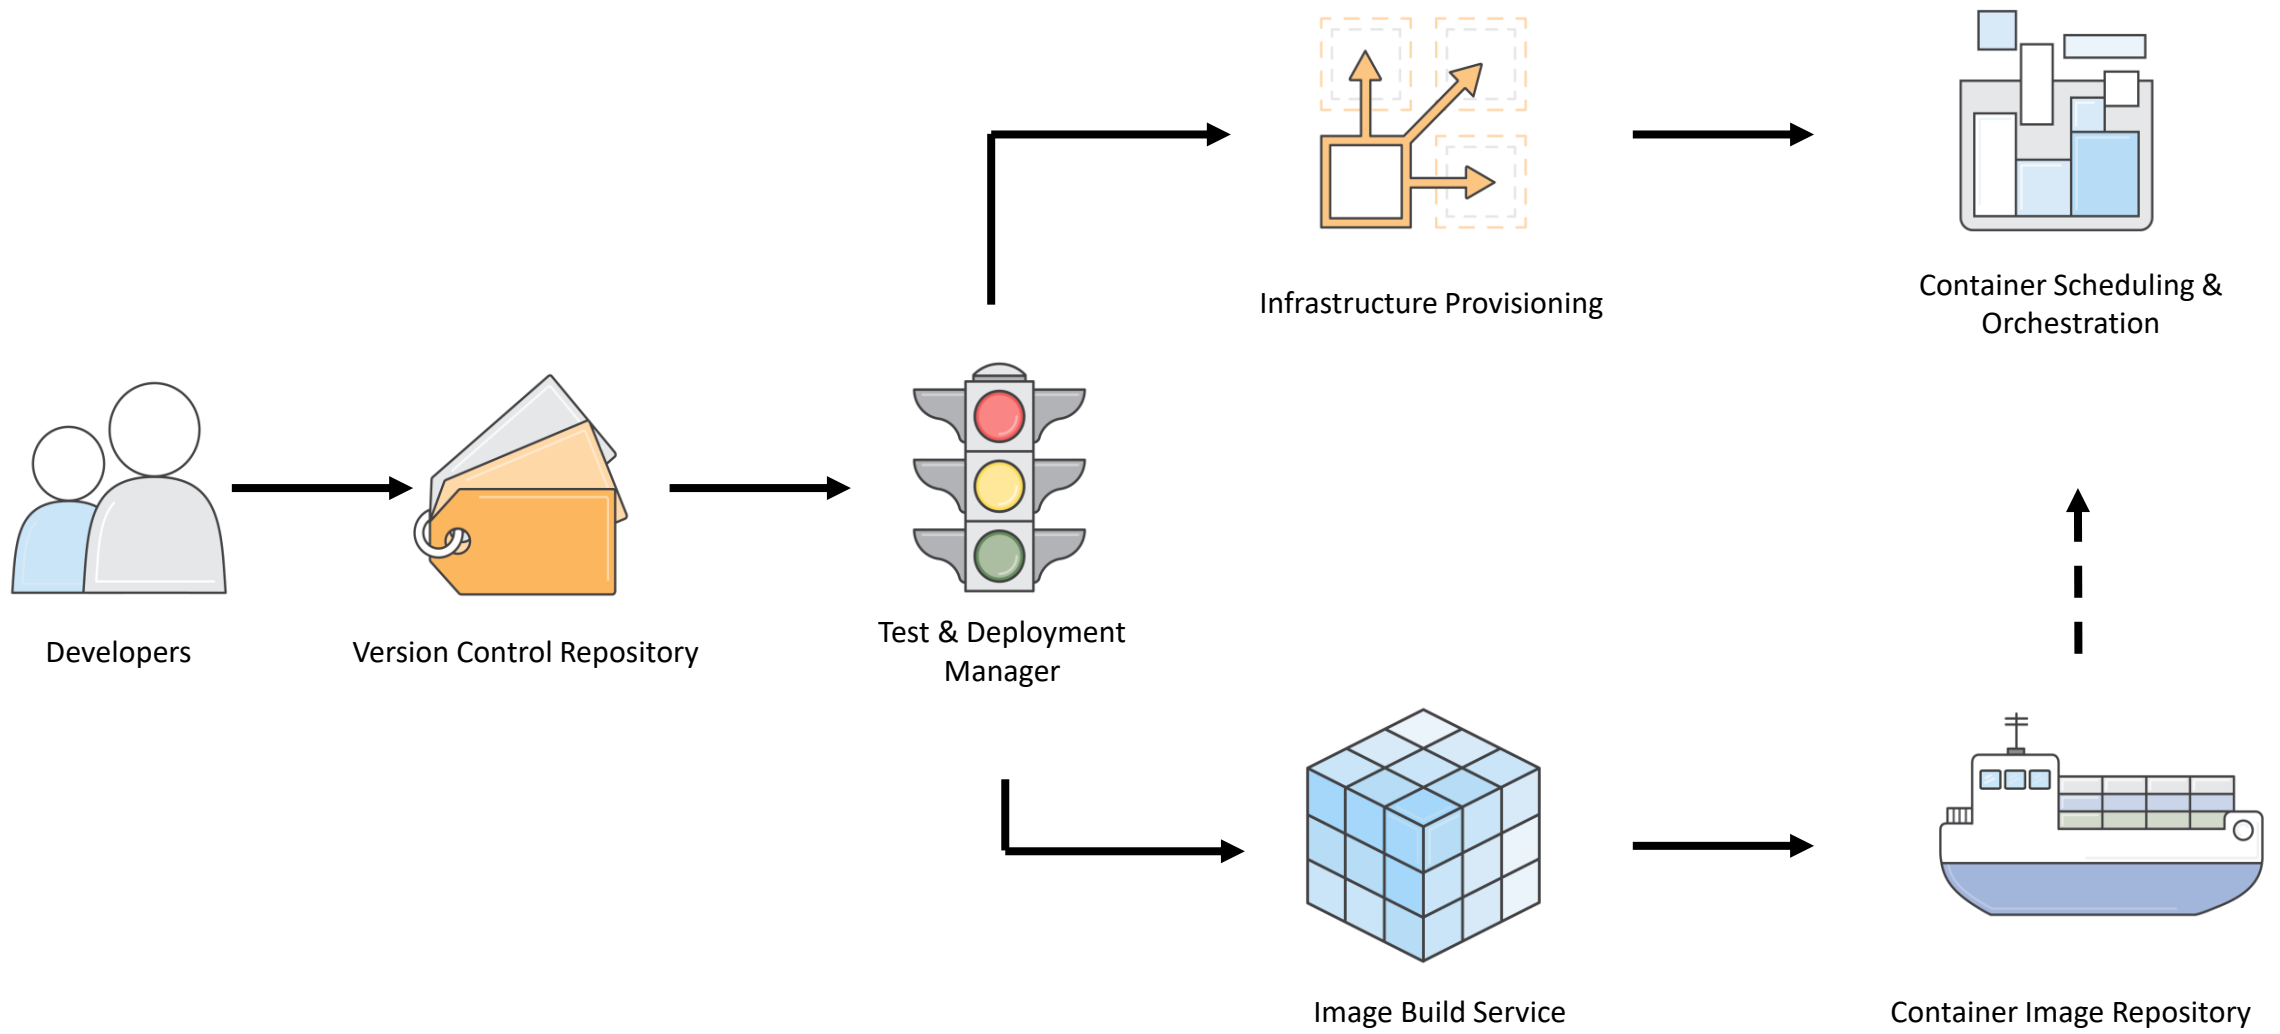
Automate the provisioning of the servers that host microservice containers:

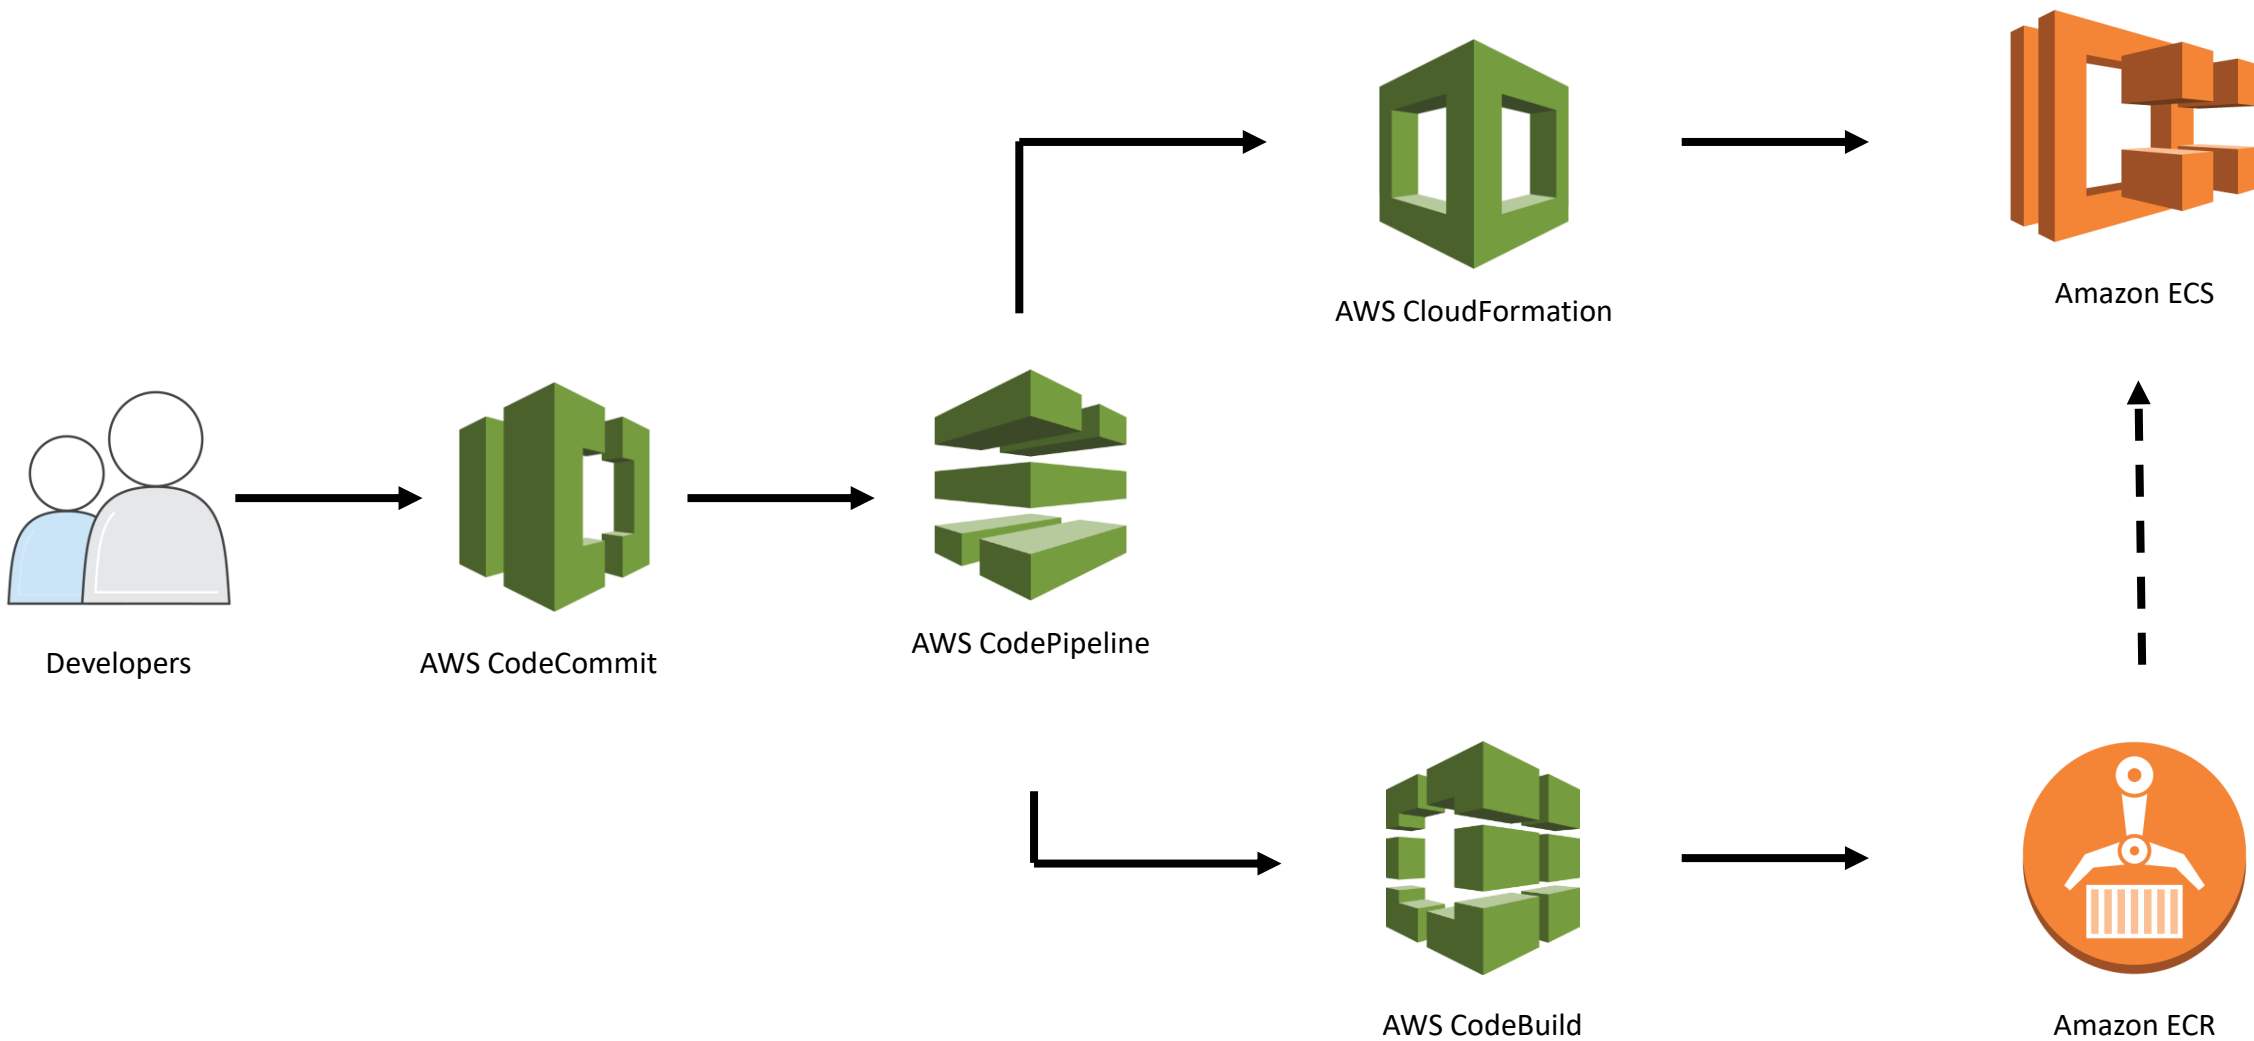
Terraform, Ansible,
Amazon CloudFormation

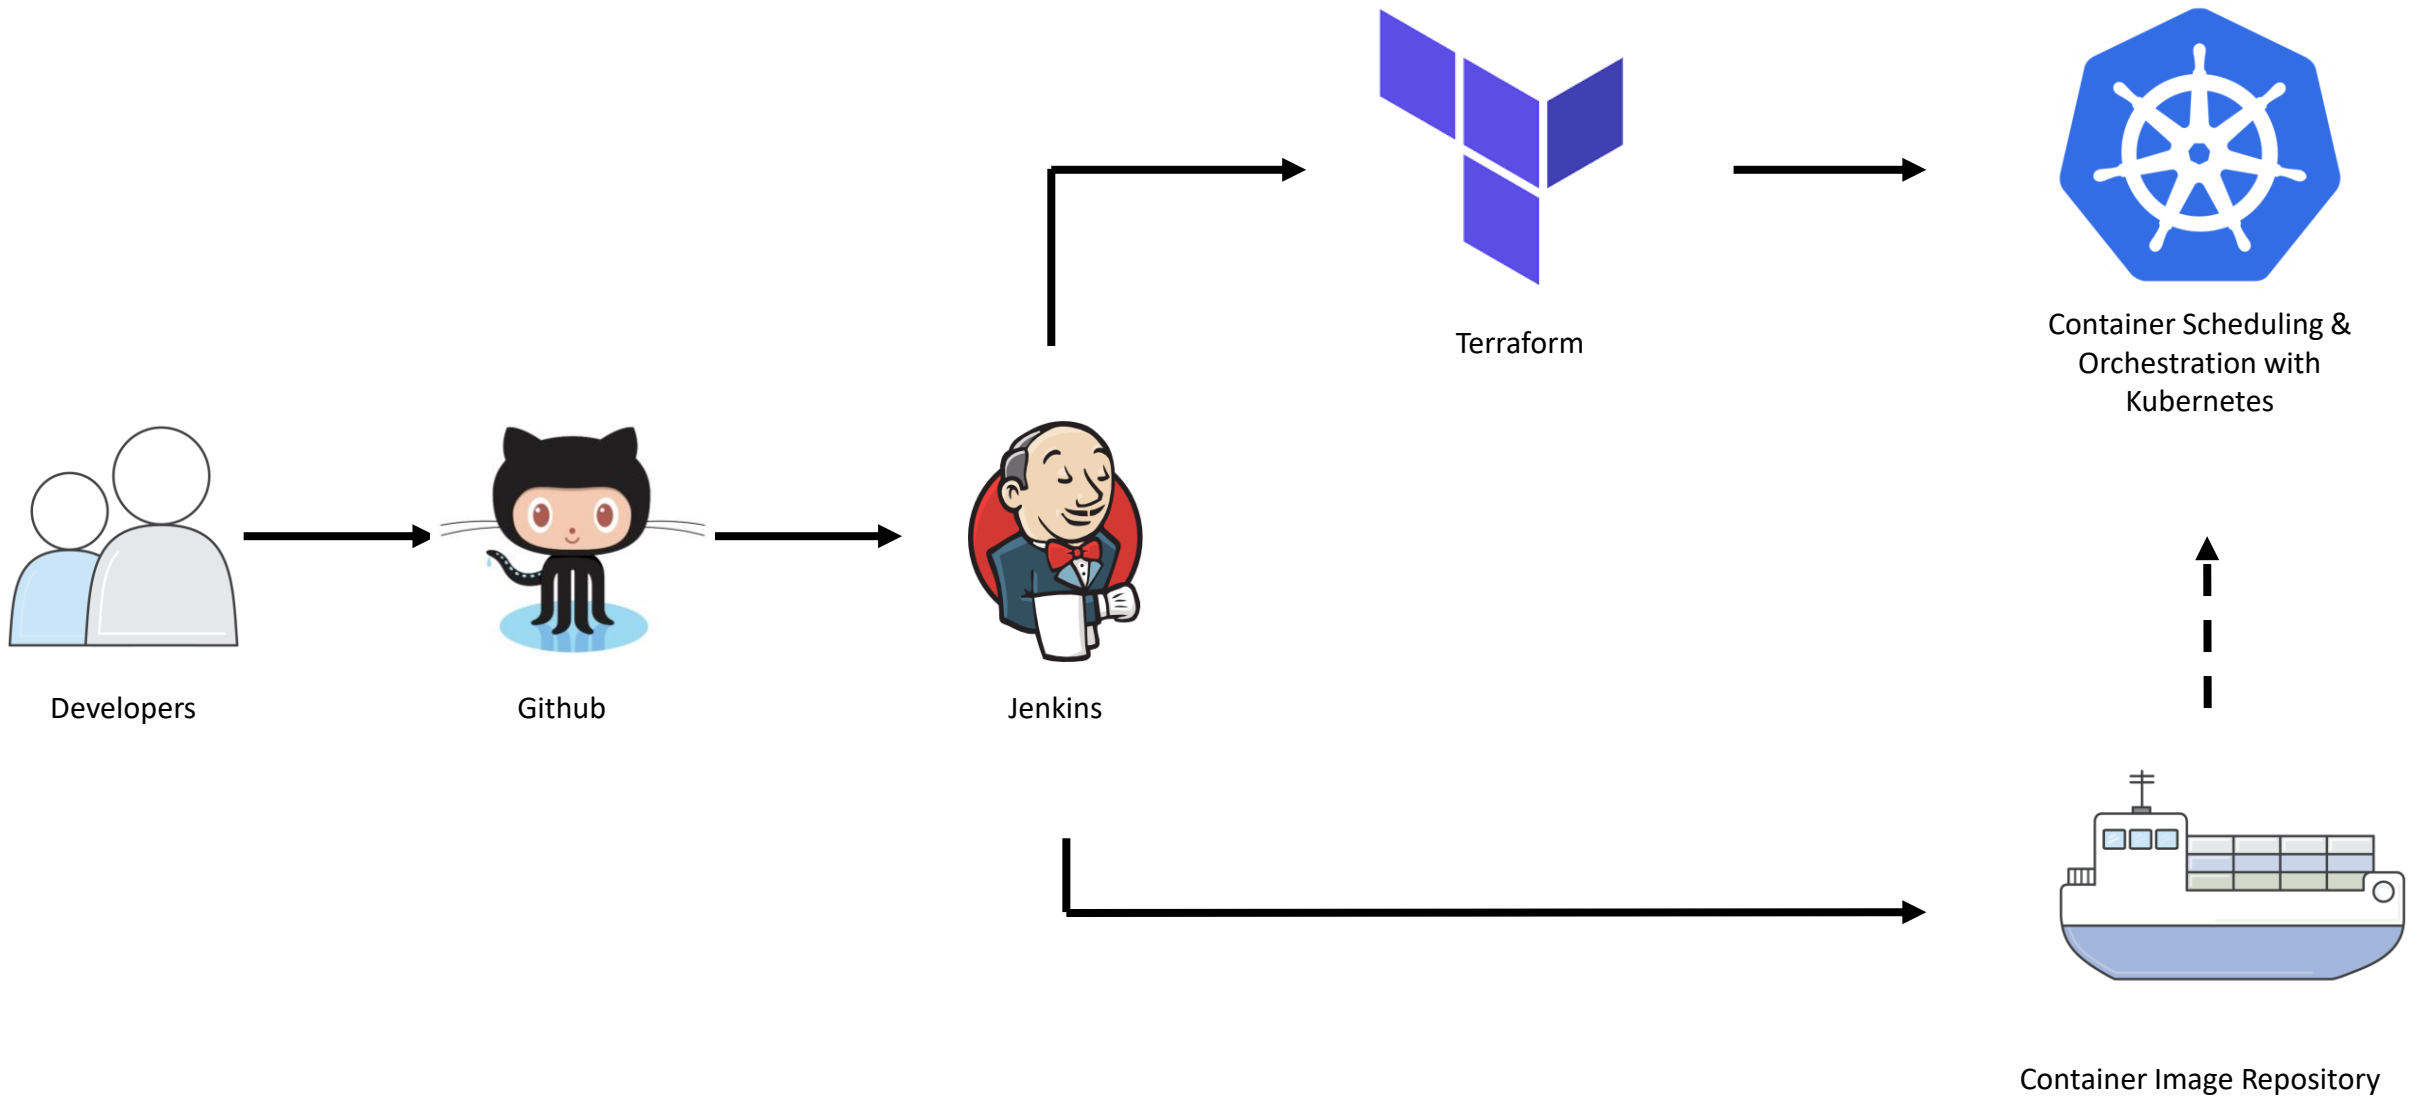


Automate the placement of
containerized service processes onto
hosts:

Amazon Elastic Container Service,
Kubernetes, Docker Swarm








Summary

12 factor application principles




Codebase




Dependencies



Configuration



Port Binding




Stateless




Fast Launch



Graceful stop




Log stream



Concurrent


Microservice principles



Automation



Componentization



Product Focused




Decentralization

Thank you!



peckn@amazon.com



nathankpeck



nathanpeck